

Code review summary

Date of Review: September 2, 2025

Reviewed By: AI Code Review Assistant (Claude)

Author of Changes: Junior Developer - Regression Model Development Team

Step 1: Analyze the codebase using AI tools

Provide a brief description of what the code update is intended to do:

This code update implements a comprehensive regression analysis pipeline for customer spending behavior prediction. The project creates both linear and logistic regression models to analyze customer spending patterns and predict high-spending behavior based on demographic and behavioral factors.

Step 1A: High-level codebase summary

This R-based machine learning project is designed as a customer analytics system for predicting spending behavior. The codebase implements a complete data science workflow including data preparation, exploratory analysis, model development, and validation.

Main Components:

- **Data Preparation Module:** Handles CSV loading, missing value detection/removal, and summary statistics calculation
- **Visualization Engine:** Creates distribution histograms and correlation scatter plots using ggplot2
- **Linear Regression System:** Develops both simple and multiple regression models for continuous spending prediction
- **Model Validation Framework:** Implements diagnostic testing including VIF analysis, normality tests, and residual analysis
- **Logistic Regression Classifier:** Predicts binary high-spending behavior with probability estimation
- **Prediction Pipeline:** Generates forecasts for new customer profiles with confidence intervals

Component Interactions: The system follows a linear workflow where cleaned data flows through exploratory visualization, then feeds into both linear and logistic modeling pipelines. The validation framework operates on both model types, and the prediction system provides actionable insights for business decision-making.

Key Dependencies:

- tidyverse (data manipulation), broom (model tidying), car (diagnostics), ggplot2 (visualization)

Step 1B: Explanation of complex function

The most complex component in this codebase is the **multiple regression model development and validation pipeline** (lines 67-95). This section combines model fitting, diagnostic testing, and assumption validation.

Purpose: Create a robust multiple regression model that predicts customer spending based on Income, Age, and PurchaseHistoryScore while ensuring statistical assumptions are met.

Logic Breakdown:

1. **Model Fitting:** Uses `lm()` to create a multiple regression with three predictors
2. **Diagnostic Analysis:** Employs `par(mfrow = c(2,2))` to create a 2x2 grid of diagnostic plots
3. **Multicollinearity Testing:** Calculates Variance Inflation Factors (VIF) to detect predictor correlation
4. **Normality Assessment:** Uses Shapiro-Wilk test and Q-Q plots to validate residual distribution
5. **Assumption Validation:** Combines visual and statistical tests to ensure model reliability

Critical Nuances:

- The VIF threshold (typically >5 or >10) determines if predictors are too correlated
- Shapiro-Wilk p-values >0.05 indicate normally distributed residuals
- The diagnostic plots reveal patterns in residuals that could violate regression assumptions

Step 1C: New code snippet

r

```
# Enhanced Model Comparison and Feature Engineering
# Added to improve model performance and interpretability

# Create interaction terms and polynomial features
analysis_df_enhanced <- analysis_df_final %>%
  mutate(
    # Interaction between Income and Age
    Income_Age_Interaction = Income * Age,
    # Polynomial term for Income (diminishing returns effect)
    Income_Squared = Income^2,
    # Age categories for better interpretation
    Age_Category = case_when(
      Age < 30 ~ "Young",
      Age >= 30 & Age < 50 ~ "Middle",
      Age >= 50 ~ "Senior"
    )
  )

# Enhanced multiple regression with feature engineering
model_enhanced <- lm(TotalSpend ~ Income + Age + PurchaseHistoryScore +
  Income_Age_Interaction + Income_Squared,
  data = analysis_df_enhanced)

# Model comparison using AIC/BIC
model_comparison <- tibble(
  Model = c("Simple", "Multiple", "Enhanced"),
  AIC = c(AIC(model_simple), AIC(model_multi), AIC(model_enhanced)),
  BIC = c(BIC(model_simple), BIC(model_multi), BIC(model_enhanced)),
  R_Squared = c(summary(model_simple)$r.squared,
    summary(model_multi)$r.squared,
    summary(model_enhanced)$r.squared)
)

print(model_comparison)
```

Commentary: This enhancement adds feature engineering capabilities to capture non-linear relationships and interaction effects between predictors. The model comparison framework allows data scientists to objectively evaluate which approach provides the best balance of explanatory power and complexity.

Step 2: Review and explain a complex function with AI assistance

Files, classes, or modules reviewed:

- Main analysis script (regression_model_final.R)
- Data input module (CSV loading section)
- Visualization components (ggplot2 implementations)
- Statistical modeling functions (lm, glm implementations)
- Diagnostic testing framework (VIF, Shapiro-Wilk, residual analysis)

Step-by-Step Explanation of Complex Function: Multiple Regression Pipeline

Step 1: Model Initialization The pipeline begins by creating a multiple regression model using `lm(TotalSpend ~ Income + Age + PurchaseHistoryScore, data = analysis_df)`. This establishes the mathematical relationship where total spending is predicted by three customer characteristics.

Step 2: Model Summary Generation The `summary(model_multi)` function generates comprehensive statistics including coefficient estimates, standard errors, t-values, and p-values. This reveals which predictors significantly influence spending behavior.

Step 3: Model Diagnostics Setup `par(mfrow = c(2,2))` creates a 2x2 plotting grid, then `plot(model_multi)` generates four diagnostic plots: residuals vs fitted values, Q-Q plot for normality, scale-location plot for homoscedasticity, and residuals vs leverage for outlier detection.

Step 4: Multicollinearity Assessment The `vif(model_multi)` function calculates Variance Inflation Factors for each predictor. VIF values above 5-10 indicate problematic correlation between predictors that could destabilize model coefficients.

Step 5: Normality Testing `shapiro.test(residuals(model_multi))` performs a formal statistical test for residual normality. Combined with Q-Q plots, this validates the assumption that model errors follow a normal distribution.

Step 6: Prediction Implementation The pipeline concludes with `predict()` functionality that generates point estimates and confidence intervals for new customer profiles, enabling practical business applications.

Step 3: Generate a new code snippet to address a feature request

Category	Yes / No / N/A	Comments
Follows style guide / naming conventions	Yes	Uses consistent snake_case naming, clear variable names, and standard R conventions
Code is modular and reusable	Partially	Functions are well-separated but could benefit from custom function definitions for repeated operations
Comments and documentation are appropriate	No	Missing inline comments explaining statistical choices and business logic
No unused variables or dead code	Yes	All variables are utilized in subsequent analysis
Logic is clear and correct	Yes	Statistical methods are properly implemented with correct syntax
Security best practices followed	N/A	Limited security concerns in analytical R code
Error handling implemented appropriately	No	Missing error handling for file loading, model convergence, and data validation
Test coverage exists / was updated	No	No unit tests or validation checks for model outputs
Performance implications considered	Partially	Efficient use of tidyverse but lacks optimization for large datasets
Backward compatibility preserved	Yes	Uses stable R packages with consistent APIs

Step 4: Update technical documentation using AI tools

Notes, feedback, and follow-up action items:

Code Quality Improvements:

- "Add comprehensive inline comments explaining statistical methodology choices"
- "Implement error handling for CSV loading and model convergence failures"
- "Consider creating custom functions for repeated model evaluation tasks"

Statistical Rigor:

- "Add model assumption testing before final model selection"
- "Include cross-validation for more robust model evaluation"
- "Consider feature scaling for models with different variable units"

Business Impact:

- "Add interpretation guidelines for stakeholders who will use model predictions"
- "Include confidence interval explanations for business decision-making"
- "Create visualization dashboard for model results presentation"

Documentation Standards:

- "Add function-level documentation following roxygen2 standards"
- "Include data dictionary explaining variable meanings and units"
- "Create README with project setup and execution instructions"

Step 5: Standardize the documentation template with AI recommendations

Brief note summarizing AI-driven improvements:

Based on AI analysis, I enhanced the documentation template to include statistical methodology explanations, business context interpretation, and reproducibility guidelines. Key improvements include adding model assumption validation checklists, prediction confidence interval documentation, and stakeholder-friendly result interpretation guides. The template now better serves both technical reviewers and business stakeholders who need to understand and act on model outputs.

- Final Review Status:** ☐ Approved
☒ Approved with minor changes
☐ Needs major revisions
☐ Rejected

Recommended Minor Changes:

- Add error handling for edge cases
- Include more comprehensive inline documentation
- Implement unit tests for critical model functions

Sign-offs

Reviewer name	Role	Signature / Initials	Date
Claude AI Assistant	Senior Code Reviewer	CAI	Sept 2, 2025
[Pending]	Technical Lead	[Pending]	[Pending]
[Pending]	Data Science Manager	[Pending]	[Pending]